

```
source <(kubectl completion bash)
```

### ● 1.Set configuration context \$kubectl config use-context k8s

Monitor the logs of Pod **foobar** and Extract log lines corresponding to error *unable-to-access-website* Write them to /opt/KULM00201/foobar

监控 foobar Pod 的日志，提取 *unable-to-access-website* 相应的行写入到/opt/KULM00201/foobar 文件中

解答：

```
#kubectl logs foobar | grep 'unable-to-access-website' > /opt/KULM00201/foobar
```

备注：

撤销 taint

```
kubectl taint node vms51.rhce.cc node-role.kubernetes.io/master-
```

### ● 2.Set configuration context \$kubectl config use-context k8s

List all PVs sorted by **name**, saving the full kubectl output to /opt/KUCC0010/my\_volumes.

Use kubectl **own** functionally for sorting the output, and do not manipulate it any further.

使用 **name** 排序列出所有的 PV，把输出内容存储到/opt/KUCC0010/my\_volumes 文件中

使用 **kubectl own** 对输出进行排序，并且不再进一步操作它。

解答：

```
#kubectl get pv --all-namespaces --sort-by={.metadata.name} > /opt/KUCC0010/my_volumes
```

### ● 3.Set configuration context \$kubectl config use-context k8s

Ensure a single instance of Pod **nginx** is running on each node of the Kubernetes cluster where **nginx** also represents the image name which has to be used. Do no override any taints currently in place.

Use **Daemonset** to complete this task and use **ds.kusc00201** as Daemonset name.

确保在 kubectl 集群的每个节点上运行一个 Nginx Pod。其中 Nginx Pod 必须使用 Nginx 镜像。不要覆盖当前环境中的任何 taints。

使用 **Daemonset** 来完成这个任务，Daemonset 的名字使用 ds.kusc00201。

引用：Concepts->Workloads->Controllers->DaemonSet

解答：

```
#cat ds.kusc00201.yaml
```

```
apiVersion: apps/v1
```

```
kind: DaemonSet
```

```
metadata:
```

```
  name: ds.kusc00201
```

```
  namespace: default
```

```
  labels:
```

```
    k8s-app: ds.kusc00201
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      name: ds.kusc00201
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        name: ds.kusc00201
```

```
    spec:
```

```
      containers:
```

```
        - name: nginx
```

```
          image: nginx
```

```
#kubectl apply -f ds.kusc00201
```

#### ●4.Set configuration context \$kubectl config use-context k8s

Perform the following tasks

Add an init container to **lumpy-koala**(which has been defined in spec file /opt/kucc00100/pod-spec-KUCC00100.yaml)

The init container should create an empty file named **/workdir/calm.txt**. If /workdir/calm.txt is not detected, the Pod should exit Once the spec file has been updated with the init container definition, the Pod should be created.

添加一个 initcontainer 到 lumpy-koala

这个 initcontainer 应该创建一个名为/workdir/calm.txt 的空文件，如果/workdir/calm.txt 没有被检测到，这个 Pod 应该更新 spec 文件并退出，这个 Pod 应该被创建

引用：Task->Configure Pod and Containers->Configure Pod Initialization(Create a Pod that has an Init Container)

Task->Configure Pod and Containers->Configure Liveness and Readiness Probes(Define a liveness command)

解答：

基础环境：

apiVersion: v1

kind: Pod

metadata:

name: myapp-pod

labels:

app: myapp

spec:

containers:

- name: myapp-container

image: nginx

volumeMounts:

- name: workdir

mountPath: /workdir

livenessProbe:

exec:

command:

- cat

- /workdir/calm.txt

initContainers:

- name: install

image: busybox

command:

- touch

- /workdir/calm.txt

volumeMounts:

- name: workdir

mountPath: /workdir

volumes:

- name: workdir

emptyDir: {}

#kubectl apply -f pod-basic.yaml

备注：

1.在外层容器上挂载目录，不然无法识别，

2.添加 livenessprobe 监测文件是否存在信息

3.添加 initContainer

#### ●5.Set configuration context \$kubectl config use-context k8s

Create a pod named **kucc4** with a single container for each of the following images running inside(there may be between 1 and 4 images specified):nginx +redis+Memcached+consul

创建一个名为 kucc4 的 Pod,其中内部运行着 nginx+redis+memcached+consul 4 个容器

引用: [Concepts->Workloads->Pods->Pod Overview](#)

找到创建 pod 的例子, 修改添加 container 的内容

解答:

```
#cat kucc4.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: kucc4
```

```
  labels:
```

```
    app: kucc4
```

```
spec:
```

```
  containers:
```

```
  - name: nginx
```

```
    image: nginx
```

```
  - name: redis
```

```
    image: redis
```

```
  - name: memcached
```

```
    image: memcached
```

```
  - name: consul
```

```
    image: consul
```

```
#kubectl apply -f kucc4.yaml
```

## ● 6. Set configuration context \$kubectl config use-context k8s

Schedule a Pod as follows:

Name: nginx-kusc00101

Image: nginx

Node selector: disk=ssd

创建 Pod, 名字为 nginx-kusc00101, 镜像为 nginx, 存放在 label 为 disk=ssd 的 node 上

引用: [Concepts->Configuration->Assigning Pods to Nodes](#)

解答:

```
#cat nginx-kusc00101.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx-kusc00101
```

```
  labels:
```

```
    env: test
```

```
spec:
```

```
  containers:
```

```
  - name: nginx
```

```
    image: nginx
```

```
  nodeSelector:
```

```
    disk: ssd
```

```
#kubectl apply -f nginx-kusc00101.yaml
```

备注:

给 node 添加 label

```
kubectl label node vms52.rhce.cc disk=ssd
```

给 node 取消 label

```
kubectl label node vms52.rhce.cc disk-
```

### ● 7. Set configuration context \$kubectl config use-context k8s

Create a deployment as follows

Name: nginx-app

Using container **nginx** with version 1.11.9-alpine

The deployment should contain **3** replicas

Next, deploy the app with new version **1.12.0**-alpine by performing a rolling update and record that update.

Finally, rollback that update to the previous version **1.11.9**-alpine.

创建 deployment

名字为 nginx-app

容器采用 1.11.9 版本的 nginx

这个 deployment 包含 3 个副本

接下来，通过滚动升级的方式更新镜像版本为 1.12.0，并记录这个更新

最后，回滚这个更新到之前的 1.11.9 版本

解答：

```
#kubectl run nginx-app --image=nginx:1.11.9 --replicas=3 --record
```

```
#kubectl set image deployment nginx-app nginx-app=nginx:1.12.0 --record
```

```
#kubectl rollout history deployment nginx-app
```

```
#kubectl rollout undo deployment nginx-app --to-revision=1
```

备注：

kubectl rollout pause deployment nginx-app 暂停 deployment，不记入 history

kubectl rollout resume deployment nginx-app 恢复 deployment 重新记入 history

### ● 8. Set configuration context \$kubectl config use-context k8s

Create and configure the service **front-end-service** so it's accessible through **NodePort/ClusterIp** and routes to the existing pod named **front-end**.

创建和配置 service，名字为 front-end-service。可以通过 NodePort/ClusterIp 开访问，并且路由到 front-end 的 Pod 上

解答：

```
#kubectl expose pod front-end --name=front-end-service --type="NodePort" --port=80
```

### ● 9. Set configuration context \$kubectl config use-context k8s

Create a Pod as follows:

Name: **jenkins**

Using image: **jenkins**

In a new Kubernetes namespace named **website-frontend**

创建一个 Pod，名字为 Jenkins，镜像使用 Jenkins。在新的 namespace website-frontend 上创建

引用：Concepts->Workloads->Pods->Pod Overview

解答：

```
#kubectl create ns website-frontend
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: jenkins
```

```
  labels:
```

```
    app: jenkins
```

```
spec:
```

```
  containers:
```

```
    - name: jenkins
```

```
      image: jenkins
```

```
#kubectl apply -f jenkins.yaml -n website-frontend
```

### ● 10.Set configuration context \$kubectl config use-context k8s

Create a deployment spec file that will:

Launch 7 replicas of the **redis** image with the **label : app\_enb\_stage=dev**

Deployment name: **kual00201**

Save a copy of this spec file to **/opt/KUAL00201/deploy\_spec.yaml** (or .json)

When you are done,clean up(delete) any new k8s API objects that you produced during this task

创建 deployment 的 spec 文件:

使用 redis 镜像, 7 个副本, label 为 app\_enb\_stage=dev

deployment 名字为 kual00201

保存这个 spec 文件到/opt/KUAL00201/deploy\_spec.yaml

完成后, 清理(删除)在此任务期间生成的任何新的 k8s API 对象

解答:

```
#kubectl run kual00201 --image=redis --replicas=7 --labels="app_enb_stage=dev" --dry-run -o yaml > /opt/KUAL00201/deploy_spec.yaml
```

### ● 11.Set configuration context \$kubectl config use-context k8s

Create a file **/opt/KUCC00302/kucc00302.txt** that lists all pods that implement Service **foo** in Namespace **production**.

The format of the file should be one pod name per line.

创建一个文件/opt/KUCC00302/kucc00302.txt , 这个文件列出所有的 service 为 foo ,在 namespace 为 production 的 Pod

这个文件的格式是每行一个 Pod

解答:

```
#kubectl get svc --show-labels -n production
```

```
#kubectl get pods -l name=lable-xxx1 -n production | grep -v NAME | awk '{print $1}' >> /opt/KUCC00302/kucc00302.txt
```

备注:

### ● 12.Set configuration context \$kubectl config use-context k8s

Create a Kubernetes Secret as follows:

Name: **super-secret**

Credential: **alice** or username: **bob**

Create a Pod named **pod-secrets-via-file** using the **redis** image which mounts a secret named **super-secret** at **/secrets**

Create a second Pod named **pod-secrets-via-env** using the **redis** image, which exports **credential/username** as

**TOPSECRET/CREDENTIALS**

备注: Concepts->Configuration->Secrets

解答:

```
# echo -n 'bob' |base64
```

```
Ym9i
```

```
#cat secret.yaml
```

```
apiVersion: v1
```

```
kind: Secret
```

```
metadata:
```

```
  name: super-secret
```

```
type: Opaque
```

```
data:
```

```
  username: Ym9i
```

```
#cat pod-secret-via-file.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod-secret-via-file
```

```
spec:
```

```

containers:
- name: redis
  image: redis
volumeMounts:
- name: foo
  mountPath: "/secret"
  readOnly: true
volumes:
- name: foo
  secret:
    secretName: super-secret
#kubectl apply -f pod-secret-via-file.yaml

```

```
#cat pod-secret-via-env.yaml
```

```

apiVersion: v1
kind: Pod
metadata:
  name: pod-secret-via-env
spec:
  containers:
  - name: redis
    image: redis
    env:
    - name: CREDENTIAL
      valueFrom:
        secretKeyRef:
          name: super-secret
          key: username

```

```
#kubectl apply -f pod-secret-via-env.yaml
```

```
#kubectl get pod
```

```
[root@vms51 cka]# kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
pod-secret-via-env	1/1	Running	0	2m
pod-secret-via-file	1/1	Running	0	11m

### ● 13.Set configuration context \$kubectl config use-context k8s

Create a pod as follows:

Name: **non-persistent-redis**

Container image: **redis**

Named-volume with name: **cache-control**

Mount path : **/data/redis**

It should launch in the **pre-prod** namespace and the volume **MUST NOT** be persistent

备注: Concept->Storage->Volumes->emptyDir(Example Pod)

解答:

```

#kubectl create ns pre-prod
#cat non-persistent-redis.yaml
apiVersion: v1
kind: Pod
metadata:
  name: non-persistent-redis
spec:

```

```
containers:
- image: redis
  name: redis-container
  volumeMounts:
  - mountPath: /data/redis
    name: cache-control
volumes:
- name: cache-control
  emptyDir: {}
#kubectl apply -f non-persistent-redis.yaml -n pre-prod
```

#### ● 14.Set configuration context \$kubectl config use-context k8s

Scale the deployment **webserver** to **6** pods

解答:

```
#kubectl scale deployment webserver --replicas=6
```

备注:

对接了 **heapster**, 和 **HPA** 联动后自动弹性伸缩

```
Kubectl autoscale deployment nginx-app --min=10 --max=15 --cpu-percent=80
```

#### ● 15.Set configuration context \$kubectl config use-context k8s

Check to see how many **nodes** are ready (not including nodes tainted NoSchedule) and write the number to **/opt/nodenum**

解答:

```
#kubectl get node | grep -i ready | wc -l
```

#### ● 16.Set configuration context \$kubectl config use-context k8s

From the Pod label **name=cpu-utilizer**, find pods running **high CPU workloads** and write the name of the Pod consuming **most CPU** to the file **/opt/cpu.txt** (which already exists)

解答:

```
#kubectl top pod -l name=cpu-utilizer
```

找到消耗 CPU 最高的 pod

```
echo 'kusc00201-5tzfk' >> /opt/cpu.txt
```

#### ● 17.Set configuration context \$kubectl config use-context k8s

Create a deployment as follows

Name: **nginx-dns**

Exposed via a service : **nginx-dns**

Ensure that the service & pod are accessible via their respective DNS records

The container(s) within any Pod(s) running as a part of this deployment should use the **nginx** image

Next, use the utility **nslookup** to look up the DNS records of the service & pod and write the output to **/opt/service.dns** and **/opt/pod.dns** respectively.

Ensure you use the **busybox:1.28** image (or earlier) for any testing, as the latest release has an upstream bug which impacts the use of **nslookup**.

创建一个 **deployment**

名字为:**nginx-dns**

路由服务名为: **nginx-dns**

确保服务和 pod 可以通过各自的 DNS 记录访问

容器使用 **nginx** 镜像

使用 **nslookup** 工具来解析 **service** 和 **pod** 的记录并写入相应的 **/opt/service.dns** 和 **/opt/pod.dns** 文件中

确保你使用 **busybox:1.28** 的镜像用来测试, 最新的版本会影响 **nslookup** 的使用

备注: Task->Administer a Cluster->Install a Network Policy Provider->Debugging DNS Resolution

解答:

```
#cat busybox1.28.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  labels:
```

```
    run: busybox28
```

```
  name: busybox28
```

```
spec:
```

```
  containers:
```

```
  - image: busybox:1.28
```

```
    name: busybox28
```

```
    command:
```

```
      - sleep
```

```
      - "3600000"
```

```
#kubectl apply -f busybox28.yaml
```

```
#kubectl run nginx-dns --image=nginx
```

```
kubectl expose deployment nginx-dns --port=80
```

```
# kubectl exec -it busybox28 -- nslookup kubernetes.default --测试
```

```
# kubectl get pod | grep nginx-dns
```

```
nginx-dns-8469864c5d-8cpjw    1/1      Running    0          58m
```

```
#kubectl exec -it busybox28 -- nslookup nginx-dns-8469864c5d-8cpjw
```

```
Server:      10.96.0.10
```

```
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local
```

```
Name:      nginx-dns-8469864c5d-8cpjw
```

```
Address 1: 101.37.252.74
```

```
# kubectl get svc | grep dns
```

```
nginx-dns    ClusterIP    10.98.24.207    <none>    80/TCP    58m
```

```
[root@vms51 ~]# kubectl exec -it busybox28 -- nslookup nginx-dns
```

```
Server:      10.96.0.10
```

```
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local
```

```
Name:      nginx-dns
```

```
Address 1: 10.98.24.207 nginx-dns.default.svc.cluster.local
```

将 nslookup 解析出来的结果写入到相应的文件中

### ● 18.No configuration context change required for this item

Create a snapshot of the etcd instance running at <https://127.0.0.1:2379> saving the snapshot to the file path

[/data/backup/etcd-snapshot.db](#)

The etcd instance is running etcd version 3.2.18

The following TLS certificates/key are supplied for connecting to the server with etcdctl

CA certificate: [/opt/KUCM00302/ca.crt](#)

Client certificate: [/opt/KUCM00302/etcd-client.crt](#)

Client key: [/opt/KUCM00302/etcd-client.key](#)

备注: Task->Administer a Cluster->Install a Network Policy Provider->Operating etcd clusters for Kubernetes

解答:

```
#export ETCDCTL_API=3
```

```
#etcdctl --endpoint=https://127.0.0.1 --cert=/opt/KUCM00302/etcd-client.crt --cacert=/opt/KUCM00302/ca.crt
```



--key=/opt/KUCM00302/etcd-client.key snapshot save /data/backup/etcd-snapshot.db

### ● 19. Set configuration context \$kubectl config use-context **ek8s**

Set the node labelled with **name=ek8s-node-1** as unavailable and reschedule all the pods running on it

解答:

```
#kubectl get node --show-labels | grep name=ek8s-node-1
```

```
#kubectl drain vms52.rhce.cc
```

### ● 20. Set configuration context \$kubectl config use-context **wk8s**

A Kubernetes worker node, labelled with **name=w8s-node-0** is in state NotReady.

Investigate why this is the case, and perform any appropriate steps to bring the node to a Ready state,

Ensuring that any changes are made permanent.

Hints:

You can ssh to the failed node using \$ssh wk8s-node-0

You can assume elevated privileges on the node with the following command \$sudo -i

解答:

是由于 kubelet 没有启动

Kubectl get node 查看一个 node 是 notReady ssh 上去

systemctl status kubelet 发现没有启动

```
#systemctl start kubelet; systemctl enable kubelet
```

### ● 21. Set configuration context \$kubectl config use-context **wk8s**

Configure the kubelet system managed service, on the node labelled with name=w8s-node-1, to Launch a Pod containing a single container of image **nginx** named **myservice** automatically. Any spec files required should be placed in the /etc/Kubernetes/manifests directory on the node.

Hints:

You can ssh to the failed node using \$ssh wk8s-node-1

You can assume elevated privileges on the node with the following command \$sudo -i

备注: Concepts->Workloads->Pods->Pod Overview(Pod Templates)

解答: 在标签为 name=w8s-node-1 的 node 节点上, 目录/etc/Kubernetes/manifests 下创建一个 yaml 文件, 文件的内容为创建一个 pod。

```
#cat myservice.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: myservice
```

```
  labels:
```

```
    app: myservice
```

```
spec:
```

```
  containers:
```

```
    - name: myservice-container
```

```
      image: nginx
```

### ● 22. Set configuration context \$kubectl config use-context **bk8s**

Given a partially-functioning Kubernetes cluster, identify symptoms of failure on the cluster.

Determine the node, the failing service and take actions to bring up the failed service and restore the health of the cluster. Ensure that any changes are made permanently.

The worker node in this cluster is labelled with **name=bk8s-node-0**

Hints:

You can ssh to the relevant nodes using \$ssh \$(NODE) where \$(NODE) is one of bk8s-master-0 or bk8s-node-0

You can assume elevated privileges on any node in the cluster with the following command

`$sudo -i`

给定一个部分功能正常的 Kubernetes cluster，在 cluster 上存在失败的迹象

确定节点、故障服务，并采取行动启动故障服务并恢复 cluster 的健康状态。确保任何更改都是永久性的。

解答：

是 kube-manager-controller 没有启动，启动就做完了

`#kubectl get cs`

能看到 controller manager 没有启动，登陆到 master 上执行以下命令

`#systemctl start kube-manager-controller.service`

### ● 23.Set configuration context \$kubectl config use-context **hk8s**

Create a persistent volume with name **app-config** of capacity **1Gi** and access mode **ReadWriteOnce**. The type of volume is **hostPath** and its location is **/srv/app-config**

备注：Concepts->Storage->Persistent Volumes(Persistent Volumes)

解答：

`#cat pv.yaml`

`apiVersion: v1`

`kind: PersistentVolume`

`metadata:`

`name: app-config`

`spec:`

`capacity:`

`storage: 1Gi`

`accessModes:`

`- ReadWriteOnce`

`hostPath:`

`path: /srv/app-config`

`#kubectl apply -f pv.yaml`